

Optimistic Temporal Difference Learning for 2048

Hung Guei , Lung-Pin Chen, *Member, IEEE*, and I-Chen Wu , *Senior Member, IEEE*

Abstract—Temporal difference (TD) learning and its variants, such as multistage TD learning and temporal coherence (TC) learning, have been successfully applied to 2048. These methods rely on the stochasticity of the environment of 2048 for exploration. In this article, we propose to employ optimistic initialization (OI) to encourage exploration for 2048, and empirically show that the learning quality is significantly improved. This approach optimistically initializes the feature weights to very large values. Since weights tend to be reduced once the states are visited, agents tend to explore those states which are unvisited or visited few times. Our experiments show that both TD and TC learning with OI significantly improve the performance. As a result, the network size required to achieve the same performance is significantly reduced. With additional tunings such as expectimax search, multistage learning, and tile-downgrading technique, our design achieves the state-of-the-art performance, namely an average score of 625 377 and a rate of 72% reaching 32 768-tiles. In addition, for sufficiently large tests, 65 536-tiles are reached at a rate of 0.02%.

Index Terms—2048, optimistic initialization (OI), reinforcement learning (RL), stochastic puzzle games, temporal difference (TD) learning, n -tuple network.

I. INTRODUCTION

The 2048 is a single-player stochastic puzzle game introduced by Cirulli [1] as a variant of *Threes!* and *1024*. This intriguing and even addictive game has been popular worldwide since it is nontrivial to master despite the simple rules [2], and has also attracted researchers to develop game-playing programs [3]. Due to its simplicity and complexity [4], 2048 is considered to be an interesting and challenging platform for evaluating the effectiveness of machine learning methods [5]–[7].

In the past, many methods were proposed for 2048. Szubert and Jaśkowski [2] applied the temporal difference (TD) learning

with n -tuple networks to 2048. In their approach, 2048-tiles were achieved at a rate of 97%. Yeh *et al.* [3] introduced the multistage TD (MS-TD) learning which improved the training by separating an entire episode into several stages. Their 3-stage TD method reached 32 768-tiles with a rate of 31.75% and even achieved the first-ever 65 536-tile. Matsuzaki [8] presented a systematic analysis on n -tuples, identified some best configurations of 8×6 -tuples and 8×7 -tuples. Jaśkowski [5] improved the performance with temporal coherence (TC) learning, which accelerated the convergence of the training by adaptively reducing the learning rate. They reached state-of-the-art (SOTA) results as follows. The 32 768-tiles reaching rate was even improved up to 70%, by using a 16-stage TC method with one second per move.

Although various methods have been continually proposed and improved, these TD methods for 2048, however, were still based on the greedy policy [9] to deterministically choose actions with maximum estimations, and thus relied on the stochastic environment to provide enough randomness for exploration. In addition, it is observed that TD learning tends not to reach large tiles when training is saturated in terms of the average score, namely when the average score increases to nearly the highest [3], [10]. This reflects a potential problem of exploration deficiency. On the other hand, recent works [3], [8], [11] tended to employ larger n -tuple networks for higher performance. Hence, as the networks become larger, the issue of exploration deficiency becomes nonnegligible. Moreover, in addition to the size of the network, the learning rate is another factor related to exploration deficiency. TC learning effectively improved the performance but converged rather fast [5], [12], possibly resulting in less exploration. Researchers [2], [5] noticed the exploration issue and tried some exploration mechanisms such as ϵ -greedy and softmax, but neither worked for 2048 according to their reports. Thus, it was simply assumed that the stochastic environment provided enough randomness, and left the efficient exploration for 2048 as an open question.

In this article, we propose to use optimistic initialization (OI) to improve the TD methods for 2048. The approach is to optimistically initialize feature weights to large values in order to encourage exploration [9], [13]. Namely, those feature weights rarely adjusted or visited tend to be high, and therefore the value adjustments are often negative, i.e., these weights tend to be reduced. Thus, agents tend to explore the less visited states next time. All feature weights eventually converge after sufficient visits.

Our experiments show that both TD and TC learning with OI significantly improve the performance. With additional tunings such as expectimax search, multistage learning, and tile-downgrading technique, our design outperforms the previous

Manuscript received 9 October 2020; revised 23 April 2021 and 5 July 2021; accepted 24 August 2021. Date of publication 3 September 2021; date of current version 15 September 2022. This work was supported in part by the Ministry of Science and Technology of the Republic of China (Taiwan) under Grant 108-2634-F-009-011, Grant 109-2634-F-009-019, and Grant 110-2634-F-009-022 through Pervasive Artificial Intelligence Research Laboratory; in part by the Higher Education Sprout Project of the National Yang Ming Chiao Tung University (NYCU) and the Ministry of Education (MOE) of Taiwan under Grant 110W265; and in part by the National Center for High-performance Computing (NCHC) of Taiwan under Grant MST108115. (*Corresponding author: I-Chen Wu.*)

Hung Guei is with the Department of Computer Science, National Yang Ming Chiao Tung University, Hsinchu 30010, Taiwan (e-mail: hguei.cs03g@nctu.edu.tw).

Lung-Pin Chen is with the Department of Computer Science, Tunghai University, Taichung 407705, Taiwan (e-mail: lbchen2007@gmail.com).

I-Chen Wu is with the Research Center for IT Innovation, Academia Sinica, Taipei 115024, Taiwan, and also with the Department of Computer Science, National Yang Ming Chiao Tung University, Hsinchu 30010, Taiwan (e-mail: icwu@csie.nctu.edu.tw).

Color versions of one or more figures in this article are available at <https://doi.org/10.1109/TG.2021.3109887>.

Digital Object Identifier 10.1109/TG.2021.3109887




Fig. 1. Segment of a gameplay episode of 2048. The player first slides the puzzle up in (a), merging some tiles and receiving 76 points to (b). After the environment generates a new tile from (b) to (c), the player continues to play by sliding the puzzle left from (c) to (d) with receiving 16 points.

SOTA results [5] and achieves new SOTA performance, namely an average score of 625 377 and a rate of 72% reaching 32 768-tiles. Even more, our method requires only 20% of network weights compared with the previous SOTA method. In addition, for sufficiently large tests, 65 536-tiles are reached at a rate of 0.02%.

The article is organized as follows. Section II reviews the rules and the related techniques for 2048. Section III introduces the OI and the optimistic methods for 2048. Section IV conducts experiments and analyses of the optimistic methods. Section V discusses potentially related techniques with directions of possible future research. Finally, Section VI concludes this article.

II. BACKGROUND

In this section, the game of 2048 as well as its related methods and techniques are introduced. Section II-A introduces the rules. Sections II-B, II-C, and II-D review the foundations and implementations of TD methods. Section II-E reviews the n -tuple networks. Section II-F reviews the tree search.

A. Rules of 2048

2048 is a single-player stochastic slide-and-merge puzzle game with the objective of sliding the puzzle to merge small tiles into large tiles to create a 2048-tile. The game is played on a puzzle with 4×4 cells, starting with two randomly placed tiles. Cells on the puzzle are either empty cells or tiles numbered with powers of 2, such as 2-tiles, 4-tiles, 65 536-tiles.

Whenever the player slides the puzzle by choosing a direction from up, down, left, and right, all tiles will be moved in the chosen direction as far as possible, i.e., until they reach either the border or another tile [2]. Illustrations are shown in Fig. 1, such as sliding up from (a) to (b), and sliding left from (c) to (d). Upon sliding the puzzle, two adjacent tiles with the same value in the chosen direction, say both v -tiles, will be merged into a single $2v$ -tile. The player receives $2v$ points as the reward of merging.

After the player slides the puzzle and merges the tiles, the environment randomly adds a new tile on an empty cell as in

Fig. 1(c) from Fig. 1(b). The newly added tile is either a 2-tile or a 4-tile, with probabilities of 0.9 or 0.1, respectively [2]. Then, the player continues to slide the puzzle, repeating the above process until there is no possible direction to move. The player wins the game if a 2048-tile is created, but the game can continue until there is no available sliding direction. The final score of a game is the total cumulative rewards of merging tiles.

B. Reinforcement Learning

RL is a machine learning method that trains an *agent* how to respond to an *environment* with an objective of maximizing the total outcome [9]. The agent continuously interacts with the environment by performing the *actions* to the current *state*, and the environment responds by providing the corresponding *rewards* and the new states.

Markov decision process (MDP) is a mathematical framework for decision-making problems, which is commonly used in RL [9]. An MDP is constructed by $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R} \rangle$, where \mathcal{S} is the finite set of states; \mathcal{A} is the finite set of actions of the state; $\mathcal{P} : \mathcal{S} \times \mathcal{A} \rightarrow \mathcal{S}$ is the state transition function; and $\mathcal{R} : \mathcal{S} \times \mathcal{A} \rightarrow \mathcal{R}$ is the immediate reward function. The MDP models problems of how an agent interacts with the environment through a sequence of actions with respect to states and rewards. An episode is a sequence of states and actions starting from the beginning till the end.

The game of 2048 can be well modeled as an MDP, in which the player is considered as an agent who takes actions to the states and receives rewards from the environment. For example, the puzzles illustrated in Fig. 1(a)–(d) can be expressed as $s_t, s'_t, s_{t+1}, s'_{t+1}$ respectively, in an episode from s_0 to s_T , as follows:

$$s_0 \cdots \rightarrow s_t \xrightarrow[r_t]{a_t} s'_t \rightarrow s_{t+1} \xrightarrow[r_{t+1}]{a_{t+1}} s'_{t+1} \rightarrow \cdots s_T. \quad (1)$$

The episode starts with an initial state $s_0 \in \mathcal{S}$. At steps t , the agent performs actions $a_t \in \mathcal{A}(s_t)$ on states $s_t \in \mathcal{S}$ to transform s_t into *afterstates* $s'_t = \psi(s_t, a_t)$, where ψ is a transition function from states to afterstates, i.e., slides the puzzle. The environment responds rewards $r_t = \mathcal{R}(s_t, a_t)$, and then changes s'_t to next states $s_{t+1} \sim \mathcal{P}(s'_t, s_{t+1})$, i.e., adds a new tile. This process repeats until a terminal state $s_T \in \mathcal{S}$ that $\mathcal{A}(s_T) = \emptyset$ is reached, where T refers to the end of the episode.

The objective of the problem of MDPs is to find a policy π that decides which action to take for any given state, and maximizes the cumulative rewards [9]. The state value function is defined as $V(s_t) = \mathbb{E}[r_t + r_{t+1} + \dots]$. Therefore, the policy $\pi : \mathcal{S} \rightarrow \mathcal{A}$ can be derived from the function as

$$\pi(s_t) = \operatorname{argmax}_{a_t} \left(r_t + \sum_{\forall s_{t+1}} \mathcal{P}(s'_t, s_{t+1}) V(s_{t+1}) \right) \quad (2)$$

where $a_t \in \mathcal{A}(s_t)$ and $s'_t = \psi(s_t, a_t)$.

C. Temporal Difference Learning

TD learning is a kind of RL method that adjusts the state estimations based in part on other learned estimations [9]. This

method has been widely applied to many game-playing programs [14]–[21], and was first applied to 2048 by Szubert and Jaśkowski [2], resulting in the first-ever TD-based program that can reach 2048-tiles.

TD(0) is the simplest form of TD learning that adjusts the estimation with only one subsequent reward and estimation [9]. After the agent performs an action and receives a reward r_t , the environment will provide the next state s_{t+1} whose value is $V(s_{t+1})$. Thus, $r_t + V(s_{t+1})$, known as the TD target, is an estimation of $V(s_t)$. Therefore, the estimation error for s_t , called the TD error, is calculated as

$$\delta_t = r_t + V(s_{t+1}) - V(s_t). \quad (3)$$

Note that in [9] $V(s_{t+1})$ is weighted by a discount factor that is disregarded for simplicity. Then, the estimation of s_t , $V(s_t)$, is adjusted with the TD error and the learning rate parameter α as

$$V(s_t) \leftarrow V(s_t) + \alpha \delta_t. \quad (4)$$

TD(0) adjusts the estimation at the current step based on the current reward and the next estimation. Its general form, n -step TD, adjusts the estimation based on n subsequent rewards and a more distant estimation [9]. Furthermore, n -step TD can be generalized to TD(λ), which adjusts the current estimation based on all subsequent estimations [9], [22]. However, n -step TD and TD(λ) are not used in this article for simplicity.

The above learning framework is to evaluate the state values, i.e., $V(s_t)$. However, from the perspective of taking actions, it is more efficient for 2048 to evaluate the afterstate values, i.e., $V(s'_t)$, instead, called the afterstate learning framework [2], [5]. With afterstate values, the policy function can be more efficient than that in (2), as described as follows:

$$\pi(s_t) = \operatorname{argmax}_{a_t} (r_t + V(s'_t)). \quad (5)$$

Similarly, the TD error is then calculated as

$$\delta_t = r_{t+1} + V(s'_{t+1}) - V(s'_t). \quad (6)$$

Depending on the order of adjusting afterstate values within an episode, forward update and backward update are both common implementations [23], [24], in which the performance of using the latter can be slightly better. In addition, Q-learning can also be employed to 2048. However, for 2048 programs with n -tuple networks, a Q-learning implementation is complicated, and its performance is significantly worse than TD learning [2], so it has not been widely used.

D. Advanced TD Methods

Multistage TD (MS-TD) learning proposed by Yeh *et al.* [3] is a kind of hierarchical TD learning [25] that divides the entire episode into multiple stages, in which each stage has an independent value function, used in [3], [5], [10], and [26]. MS-TD learning improves the performance at the cost of additional storage for stages. The work in [3] applied this method with 3-stage to 2048 and obtained the first-ever 65 536-tile.

TC learning proposed for 2048 by Jaśkowski [5] is a TD variant with adaptive learning rates [12]. Instead of adjusting the

learning rate α directly, this method introduces new parameters β_i for the i th feature weight, denoted by θ_i , to modulate the adjustments as

$$\theta_i \leftarrow \theta_i + \alpha \beta_i \delta_t. \quad (7)$$

β_i represents the coherence of θ_i , and is calculated from two parameters E_i and A_i for each weight, as

$$\beta_i = \begin{cases} |E_i|/A_i, & \text{if } A_i \neq 0 \\ 1, & \text{otherwise.} \end{cases} \quad (8)$$

Both E_i and A_i are initialized with 0 and adjusted by

$$E_i \leftarrow E_i + \delta_t \text{ and } A_i \leftarrow A_i + |\delta_t|. \quad (9)$$

Therefore, the amount of adjustment is automatically reduced by coherence β_i . TC learning is an effective learning rate decay method with an overhead of triple required memory. This technique can be integrated with MS-TD, e.g., a 16-stage TC method was able to achieve 609 104 points and a 70% chance of reaching 32 768-tiles on average [5].

E. N -Tuple Networks

A straightforward method to estimate state values $V(s)$ is to use a tabular implementation for the whole state space. However, the state space requires $18^{(4 \times 4)}$ for 2048, too large to be implemented. Therefore, a function approximator is applied in practice. N -tuple network is a function approximator that has been successfully applied to applications such as Connect4 [20], Othello [21], pattern recognition [27], as well as 2048 [2], [3], [5], [8].

For n -tuple networks for 2048, an n -tuple ϕ is a sequence of distinct features, each representing a tuple of n designated cells on the puzzle. For example, let ϕ_{R1} be a 4-tuple that denotes features from the first row. For a puzzle s , say the one in Fig. 1(a), a feature $\phi_{R1}(s)$ refers to (32, 2, 0, 0). Similarly, for s' , say the one in Fig. 1(b), a feature $\phi_{R1}(s')$ refers to (64, 4, 16, 2).

An n -tuple network is an implementation for a set of weights of n -tuple features. In order to access these feature weights, let ϕ be an n -tuple associated with a lookup table (LUT) in which the feature weight of $\phi(s)$ is stored at a distinct $\text{LUT}[\phi(s)]$. An illustration of the implementation for ϕ_{R1} is as follows. The lookup table LUT_{R1} for ϕ_{R1} needs to contain c^4 distinct feature weights, since ϕ_{R1} consists of 4 cells, each with c distinct cell values. Intrinsically, c is 18, i.e., from empty cell to 131 072-tile, but is usually set to 16 or 17 for efficiency, since 65 536-tile and 131 072-tile are rarely obtained.

In this article, we define an $m \times n$ -tuple network to be m different n -tuples ϕ_1, \dots, ϕ_m with their corresponding lookup tables $\text{LUT}_1, \dots, \text{LUT}_m$. Given a state s , the state value estimation $V(s)$, is calculated by summing all of the m feature weights $\text{LUT}_i[\phi_i(s)]$ of state s as

$$V(s) = \sum_{i=1}^m \text{LUT}_i[\phi_i(s)]. \quad (10)$$

When adjusting a state estimation $V(s)$ by a TD error δ , the adjustment is equally distributed to m feature weights of state s . The following equation shows how δ is distributed to a feature




Fig. 2. (a), (b), (d), and (e) 4×6 -tuple network proposed by Yeh *et al.* [3]. (a)–(e) 5×6 -tuple network used by Jaśkowski [5].

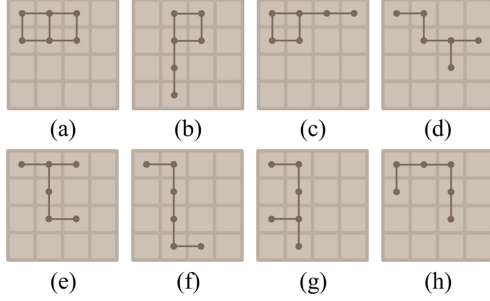


Fig. 3. (a)–(h) 8×6 -tuple network proposed by Matsuzaki [8].

weight $\text{LUT}_i[\phi_i(s)]$ in terms of TD(0) with learning rate α

$$\text{LUT}_i[\phi_i(s)] \leftarrow \text{LUT}_i[\phi_i(s)] + (\alpha/m)\delta. \quad (11)$$

Note that only the m feature weights corresponding to state s need to be adjusted, with the same adjustment $(\alpha/m)\delta$.

For 2048, symmetric sampling is a widely used technique that shares feature weights of tuples eight times by rotating and mirroring [2], [3], [5], [8]. A symmetrically sampled $m \times n$ -tuple network involves $8m$ n -tuples actually, which improves the overall performance without additional lookup tables. For example, let $\phi_{R1'}$ be a 4-tuple produced by rotating ϕ_{R1} counterclockwise, then, its feature $\phi_{R1'}(s)$ is (16, 8, 32, 32). Both ϕ_{R1} and $\phi_{R1'}$ share the same lookup table LUT_{R1} . Note that an $m \times n$ -tuple network refers to a symmetrically sampled $m \times n$ -tuple network in the rest of this article for simplicity.

Designing an effective $m \times n$ -tuple network is not trivial, and has been investigated as a research topic for 2048 [2], [3], [8], [11]. Fig. 2(a),(b),(d), and (e) illustrates the 4×6 -tuple network proposed by Yeh *et al.* [3], [10]; Fig. 2(a)–(e) illustrates the 5×6 -tuple network used by Jaśkowski [5]. Fig. 3 illustrates a full set of 6-tuples proposed by Matsuzaki [8], where the best $k \times 6$ -tuple network consists of the first k listed 6-tuples, e.g., the 8×6 -tuple network contains all from Fig. 3(a)–(h).

F. Expectimax Search

Expectimax search is a technique for stochastic games whose game tree is composed of max nodes and chance nodes [28]–[30], corresponding to states and afterstates respectively. Namely, a max node of state s_t searches all its afterstates s'_t to find the best action; and a chance node of afterstate s'_t is evaluated based on either the expected value of next states, or the TD value $V(s'_t)$ when the ply limit is reached. For 2048 with afterstate learning framework, a p -ply fixed-depth search tree has at most p layers of chance nodes.

To avoid redundant search, the search is usually integrated with a transposition table, which caches the previously seen

states and associated values. A transposition table for 2048 is implemented using hashing techniques such as *Zobrist hashing* [31], [32] and *MurmurHash* [33], [34].

III. OPTIMISTIC METHODS

In this section, optimistic methods for 2048 are introduced. Section III-A reviews the potential problems of insufficient exploration in previous works. Sections III-B and III-C present TD learning with OI for 2048. Finally, Section III-D describes how to determine initial values and use them in n -tuple networks.

A. Insufficient Exploration

In RL, the exploration–exploitation dilemma has been intensively studied by researchers for decades [9], [13], [35]. Many methods, including upper confidence bound (UCB), ϵ -greedy, softmax, as well as OI, are proposed to balance between exploration and exploitation [36] for better learning performance.

Most previous RL related works for 2048 [2], [3], [5], [6], [8], [10], [11], [23], [26] are based on simple TD method or some of its variant. Among these works, the agents simply follow the greedy policy with respect to the estimations, i.e., always select an action with maximum value during training. Past works [2], [5] have already noticed this issue, and have tried some explicit exploration techniques, such as ϵ -greedy and softmax. However, they did not succeed to have these techniques work for 2048. As a result, these TD methods involve no explicit exploration to choose actions and fully rely on the stochasticity of the environment for exploration.

For example, MS-TD learning is proposed to cope with the issue that TD learning tends not to reach large tiles [3], [10]. MS-TD learning improves performance by employing additional networks, while the issue of potential exploration deficiency of each stage remains not addressed. Another effective technique is TC learning, which reduces weight adjustments to automatically accelerate network convergence [5]. This approach is not for improving exploration, but for improving exploitation with fast convergence. However, fast convergence exacerbates the exploration issue, especially when applying a large n -tuple network.

B. Optimistic Initialization (OI)

OI is an approach that employs optimistic initial estimations to encourage exploration [9], [13]. This approach has been widely applied to RL applications, and is considered to have good convergence in practice [9]. Instead of setting the estimations to zero or random, the technique optimistically initializes them to a large value to encourage the agent to explore. Due to the large value, the estimations tend to be reduced once the corresponding states are visited, therefore leading an agent to select unexplored or rarely explored actions the next time when it revisits the same state. This process repeats until all the actions are sufficiently explored, even if the greedy policy is always applied during the training.

The estimations will eventually converge and may even converge to a near-optimal policy when the initial value is set sufficiently large [35], [37], namely the upper bound of the

value function. Thus, the learning algorithm needs to explore unvisited or rarely visited states and reduce the corresponding estimations before exploiting the best one [38]. Therefore, OI may significantly increase the training time in exchange for encouraging exploration.

C. Optimistic TD Methods

Past research summarized that nongreedy behaviors, e.g., ε -greedy, significantly inhibit the learning performance [2], [5]. In this article, we propose to use OI as an exploration mechanism for 2048. The proposed OI methods perform exploration while conserving the greedy behavior, therefore mitigating the previous inhibition phenomenon. In addition, since the explicit exploration technique is independent of some existing learning methods such as multistage and TC learning, OI can be easily applied together with these methods.

In this article, the first objective is to improve the existing TD and TC methods for 2048 by using OI, which forms the optimistic TD (OTD) and the optimistic TC (OTC) learning, respectively. Second, we propose a two-phase optimistic method, called OTD+TC learning, a hybrid learning paradigm that combines the advantages of both methods. OTD+TC learning first performs TD learning with a fixed learning rate to further encourage exploration for a while, and then, in the second phase, continues with TC fine-tuning for exploitation.

When compared to TD or TC, OTD+TC learning includes two new hyperparameters, V_{init} and P_{TC} : V_{init} is the initial value of the function approximator; and P_{TC} is the proportion of TC fine-tuning phase to the total phases. If V_{init} is set to 0, OTD+TC becomes nonoptimistic. If P_{TC} is set to 0% (or 100%), OTD+TC becomes pure OTD (or OTC).

D. Initial Values for OI

Based on the proof of the optimistic Q-learning in [37], the initial value V_{init} should be set to the theoretical maximum to ensure that the network converges to a near-optimal policy. However, the theoretical maximum for 2048 is an extremely large number that is nearly impossible to obtain. Using such a large V_{init} wastes too much time, it is nontrivial to choose V_{init} such that training result and training time are balanced.

In this article, the initial value V_{init} is chosen as illustrated as follows. Consider using an $m \times n$ -tuple network for training with nonoptimistic TD learning, and set V_{init} an estimated average score. To initialize the network feature weights, we evenly distribute the value V_{init} over these feature weights as

$$\text{LUT}[i] \leftarrow V_{\text{init}}/m, \quad \text{for all } i. \quad (12)$$

IV. EXPERIMENTS

In this section, experiments are presented to analyze the effectiveness of optimistic TD learning for 2048. Common training settings are described as follows. For convergence, each n -tuple network was trained with 100 million episodes by using the afterstate learning framework. The learning performance was evaluated every 1 million training episodes, and each performance evaluation consists of 100 000 testing episodes. For statistics, each method was trained with five runs, each with

TABLE I
PERFORMANCE OF OTD METHODS IN THE 4×6 -TUPLE NETWORK

V_{init}	Average Score	8192 [%]	16 384 [%]	32 768 [%]
0	251 920 \pm 11 217	92.87 \pm 0.90%	65.23 \pm 2.86%	0.00 \pm 0.00%
40k	254 784 \pm 13 034	93.20 \pm 1.08%	65.89 \pm 3.60%	0.00 \pm 0.00%
80k	253 393 \pm 11 732	92.52 \pm 1.13%	65.25 \pm 3.35%	0.00 \pm 0.00%
160k	244 875 \pm 10 590	91.40 \pm 1.41%	62.29 \pm 3.27%	0.00 \pm 0.00%
320k	239 075 \pm 14 781	89.68 \pm 2.11%	59.73 \pm 5.00%	0.00 \pm 0.00%
640k	228 032 \pm 1056	87.64 \pm 0.44%	55.63 \pm 0.48%	0.00 \pm 0.00%




Fig. 4. Average scores of OTC methods in the 4×6 -tuple network.

one individual trained network with a different initial random seed. The experiments were performed on workstations with Intel Xeon E5 processors. To speed up the training process, 20 threads lock-free optimistic parallelism was applied as in [5]. Specific settings will be described in each experiment below.

The results are presented in tables and figures in which each value represents the average of five trained networks. For all tables (e.g., Table I), a value indicates the average result after 100 million training episodes and is accomplished with the 95% confidence interval. For all figures (e.g., Fig. 4), a point at time t indicates the average result after t million training episodes.

The experiments are organized as follows. Section IV-A analyzes the effectiveness of initial values on OTD and OTC learning. Section IV-B analyzes OTD+TC learning with different fine-tuning proportions on networks of different sizes. Section IV-C further improves the performance by fine-tuning with other techniques for OTD+TC learning.

A. OTD and OTC Learning

In order to demonstrate the effectiveness of OI and obtain an appropriate initial value, we analyze how OTD and OTC learning perform with various initial values. Two different n -tuple networks, Yeh's 4×6 -tuple (see Fig. 2) and Matsuzaki's 8×6 -tuple (see Fig. 3), were used. Based on their average scores with nonoptimistic TD learning, we test the initial value V_{init} to 0, 40 000, 80 000, 160 000, 320 000, 640 000, respectively. For OTD learning, the learning rate α was initially set to 0.1 and was reduced to 0.01 and 0.001 after completing 50% and 75% of training; for OTC learning, a larger initial learning rate is required, so α was set to 1.0 without manual decay as in [5]. The results are organized as follows. OTD and OTC learning for the 4×6 -tuple will be discussed first, and then followed by OTD and OTC learning for the 8×6 -tuple.

For the 4×6 -tuple network, OTD learning did not obviously improve the learning quality, as given in Table I. On the other

TABLE II
PERFORMANCE OF OTC METHODS IN THE 4×6 -TUPLE NETWORK

V_{init}	Average Score	8192 [%]	16 384 [%]	32 768 [%]
0	265 481 \pm 3318	93.36 \pm 1.07%	68.33 \pm 0.52%	0.00 \pm 0.00%
40k	263 284 \pm 2445	93.68 \pm 0.32%	66.72 \pm 1.79%	0.00 \pm 0.00%
80k	265 703 \pm 2471	93.60 \pm 0.26%	67.93 \pm 1.49%	0.00 \pm 0.00%
160k	264 058 \pm 1412	93.24 \pm 0.28%	67.47 \pm 0.35%	0.00 \pm 0.00%
320k	280 281 \pm 5267	93.11 \pm 0.53%	67.83 \pm 1.50%	5.30 \pm 0.58%
640k	252 254 \pm 6836	89.18 \pm 1.48%	60.22 \pm 1.98%	3.25 \pm 0.38%




Fig. 5. Maximum scores of OTC methods in the 4×6 -tuple network.

TABLE III
PERFORMANCE OF OTD METHODS IN THE 8×6 -TUPLE NETWORK

V_{init}	Average Score	8192 [%]	16 384 [%]	32 768 [%]
0	309 208 \pm 5788	97.24 \pm 0.58%	85.13 \pm 0.78%	0.00 \pm 0.00%
40k	306 438 \pm 2400	97.27 \pm 0.11%	84.11 \pm 0.44%	0.00 \pm 0.00%
80k	365 364 \pm 4856	97.16 \pm 0.17%	85.59 \pm 0.67%	21.23 \pm 1.04%
160k	369 172 \pm 3341	97.18 \pm 0.30%	85.70 \pm 0.63%	22.47 \pm 1.31%
320k	361 471 \pm 5473	96.77 \pm 0.28%	84.41 \pm 0.96%	21.75 \pm 0.74%
640k	339 492 \pm 6747	95.79 \pm 0.28%	81.80 \pm 1.07%	18.11 \pm 2.12%

hand, OTC learning significantly improved the performance, especially for the reaching rate of 32 768-tiles. As the results given in Table II and Fig. 4, OTC with 320 000 outperformed nonoptimistic TC ($V_{\text{init}} = 0$), and even achieved a rate of 5.30% reaching 32 768-tiles, which is the first-ever reaching rate with such a 1-stage network. As shown in Fig. 5, many methods reached maximum scores around 385 000. This implicitly indicates to achieve states with 16 384-tile, 8192-tile, 4096-tile, and 2048-tile. Note that a method can continue to obtain much more scores if a 32 768-tile is reached. Only OTC with 320 k and 640 k provided enough exploration to achieve 32 760-tiles. Interestingly, as shown in Fig. 4, OTC with 640 k reached 32 768-tiles earlier than OTC with 320 000, but eventually ended at a worse average score. In general, a large V_{init} needs more time to converge. Our conjecture is that 640 000 is too large to converge. Thus, we prefer 320 000 for obtaining higher average scores.

For the 8×6 -tuple network, in contrast to the 4×6 -tuple, OTD learning significantly improves the performance. As given in Table III and Fig. 6, networks initialized with appropriate values ($80\,000 \leq V_{\text{init}} < 640\,000$) achieved higher scores, and even reached 32 768-tiles well, especially for $V_{\text{init}} = 160\,000$ and 320 000. As shown in Fig. 7, networks trained with low exploration ($V_{\text{init}} < 80\,000$) still stuck at the barrier of 32 768-tiles; and networks trained with 640 000 still did not




Fig. 6. Average scores of OTD methods in the 8×6 -tuple network.




Fig. 7. Maximum scores of OTD methods in the 8×6 -tuple network.

TABLE IV
PERFORMANCE OF OTC METHODS IN THE 8×6 -TUPLE NETWORK

V_{init}	Average Score	8192 [%]	16 384 [%]	32 768 [%]
0	310 259 \pm 4056	96.54 \pm 0.13%	84.81 \pm 0.43%	0.00 \pm 0.00%
40k	310 347 \pm 6201	96.72 \pm 0.71%	83.59 \pm 1.38%	0.00 \pm 0.00%
80k	311 214 \pm 2205	97.07 \pm 0.22%	83.45 \pm 0.23%	0.00 \pm 0.00%
160k	361 298 \pm 4833	96.34 \pm 0.20%	84.12 \pm 0.52%	22.26 \pm 2.12%
320k	360 228 \pm 18 829	95.73 \pm 0.39%	82.28 \pm 2.25%	22.74 \pm 2.38%
640k	325 357 \pm 33 416	92.76 \pm 2.55%	75.88 \pm 6.10%	17.71 \pm 4.89%

converge. On the other hand, the results of OTC learning are given in Table IV. OTC learning also performed well, but its performance was slightly worse than OTD for the 8×6 -tuple network. Interestingly, for $V_{\text{init}} = 80\,000$ on the 8×6 -tuple network, OTD achieved 32 768-tiles, while OTC did not. Since such a phenomenon did not occur on the 4×6 -tuple, it can be derived from this observation that exploration may be restricted by TC method since the weight adjustments drop too fast, therefore, a pure TD method plays an important role for larger networks.

B. OTD+TC Learning

This experiment analyzes OTD+TC learning. In order to analyze the correlation between OI and network size in detail, five experimental network sizes were chosen, and were divided into two classes: the smaller networks, including Yeh's 4×6 -tuple and 5×6 -tuple (see Fig. 2); and the larger networks, including Matsuzaki's 6×6 -tuple, 7×6 -tuple, and 8×6 -tuple (see Fig. 3). Since 320 000 performed well in most cases as described above, V_{init} is set to 320 000 in the rest of the experiments.

First, we would like to analyze the effectiveness of OTD+TC learning. The average scores and the 32 768-tile reaching rates

TABLE V
DIFFERENCE BETWEEN OTD, OTC, AND OTD+TC METHODS IN N-TUPLE NETWORKS OF DIFFERENT SIZES

Network	OTD	OTD+TC	OTC
4×6-tuple	239 075 ± 14 781 (0.00 ± 0.00%)	261 433 ± 3143 (0.02 ± 0.02%)	280 281 ± 5267 (5.30 ± 0.58%)
5×6-tuple	272 904 ± 7788 (0.01 ± 0.01%)	279 653 ± 8707 (0.47 ± 1.08%)	313 948 ± 2433 (11.48 ± 0.97%)
6×6-tuple	314 324 ± 8883 (8.33 ± 5.58%)	337 951 ± 4655 (14.39 ± 0.96%)	324 714 ± 17 688 (12.88 ± 5.88%)
7×6-tuple	352 986 ± 3553 (18.43 ± 0.69%)	360 286 ± 2332 (19.18 ± 0.70%)	335 042 ± 23 023 (16.56 ± 6.61%)
8×6-tuple	361 471 ± 5473 (21.75 ± 0.74%)	370 907 ± 1630 (22.18 ± 0.53%)	360 228 ± 18 829 (22.74 ± 2.38%)

Networks initialized with $V_{\text{init}} = 320$ k; OTD+TC used $P_{\text{TC}} = 10\%$.

TABLE VI
PERFORMANCE OF OTD+TC METHODS IN THE 8 × 6-TUPLE NETWORK

P_{TC}	Average Score	8192 [%]	16 384 [%]	32 768 [%]
10%	370 907 ± 1630	97.26 ± 0.16%	85.43 ± 0.34%	22.18 ± 0.53%
20%	371 198 ± 8301	97.19 ± 0.53%	85.55 ± 1.21%	22.78 ± 1.39%
30%	366 868 ± 7767	97.07 ± 0.45%	84.82 ± 1.28%	22.28 ± 1.38%
50%	355 806 ± 13 994	96.32 ± 0.73%	83.64 ± 1.97%	20.61 ± 2.24%
70%	344 570 ± 20 316	95.99 ± 1.06%	82.04 ± 3.58%	17.78 ± 2.83%
90%	298 673 ± 18 995	95.25 ± 1.64%	80.05 ± 4.60%	0.86 ± 2.00%

Networks were trained with initial value $V_{\text{init}} = 320$ k.




Fig. 8. Average scores of OTD+TC methods in the 8 × 6-tuple network.

of OTD, OTC, and OTD+TC with $P_{\text{TC}} = 10\%$ in the five networks are given in Table V. OTD+TC outperforms OTD in all cases, and slightly outperforms OTC in most cases of using larger networks. In the cases of using smaller networks, OTC outperforms OTD+TC; while OTD+TC hardly obtained 32 768-tiles. In brief, OTC learning performed well regardless of the network size, and OTD+TC learning outperforms OTC learning in most larger networks.

Second, it is interesting to investigate how the fine-tuning proportions P_{TC} affects the performance. Table VI gives the results of more P_{TC} of 10%, 20%, 30%, 50%, 70%, and 90% for the 8 × 6-tuple network. Figs. 8 and 9 illustrate the learning curves of 10%, 50%, 70%, and 90%, with also OTD and OTC as baselines. Based on the results, it is observed that high proportions (e.g., 70% and 90%) limit the ability to explore, and lead to worse results. Therefore, OTD+TC learning should be performed with low fine-tuning proportions such as 10% or 20% to provide sufficient exploration, as well as prevent potential overfitting issue in practice.




Fig. 9. Maximum scores of OTD+TC methods in the 8 × 6-tuple network.

TABLE VII
PERFORMANCE OF OPTIMISTIC TD METHODS TOGETHER WITH EXPECTIMAX SEARCH IN THE 4 × 6-TUPLE AND THE 8 × 6-TUPLE NETWORKS

Network	Search	Average Score	8192 [%]	16 384 [%]	32 768 [%]
4×6-tuple	1-ply	280 123 ± 4920	93.00 ± 0.61%	67.79 ± 1.39%	5.32 ± 0.43%
	3-ply	417 712 ± 2988	99.66 ± 0.10%	94.95 ± 0.60%	37.18 ± 1.62%
	5-ply	445 085 ± 8843	99.80 ± 0.89%	96.60 ± 5.02%	51.20 ± 3.58%
8×6-tuple	1-ply	370 194 ± 4366	97.23 ± 0.18%	85.30 ± 0.46%	22.18 ± 0.92%
	3-ply	475 126 ± 6407	99.73 ± 0.08%	96.88 ± 0.78%	50.73 ± 1.56%
	5-ply	500 098 ± 8590	100.00 ± 0.00%	98.20 ± 3.85%	57.80 ± 2.97%

The 4 × 6-tuple and the 8 × 6-tuple were trained with OTC and OTD+TC, respectively.

C. Further Improvements

In order to further improve the performance, we demonstrate how to apply optimistic methods together with other established techniques, such as expectimax search, multistage learning, and tile-downgrading.

1) *Expectimax Search*: As also described in [2] and [5], expectimax search can be used to further improve the performance. We applied expectimax search to two cases, the 4 × 6-tuple network with OTC and the 8 × 6-tuple with OTD+TC, since both performed generally well as shown above. The performance of 3 different depths, 1-ply, 3-ply, and 5-ply, were evaluated for the tests of 1 million, 10 000, and 100 episodes, respectively. Each evaluation is also repeated five times by using the five trained networks. Transposition tables were used to speed up the search.

Table VII gives the performance of the 4 × 6-tuple and the 8 × 6-tuple network with 1-ply, 3-ply, and 5-ply expectimax search. Benefited from OI, both average scores and 32 768-tile reaching rates achieved SOTA with respect to the same 1-stage networks and search depths, to the best of our knowledge. However, the performance was saturated around 5-ply.

2) *Multistage Learning*: We further evaluate optimistic TD learning together with multistage paradigm [3], named multistage OTD+TC (MS-OTD+TC) learning. As in Section IV-C1, we also use the two cases, the 4 × 6-tuple network with OTC and the 8 × 6-tuple network with OTD+TC. For both, the first stages started with initial states, and the second stages started with states with 16 384-tile. Note that the first stages are directly copied from trained networks, and the second stages also received 100 million training episodes. All stages were trained

TABLE VIII
PERFORMANCE OF MULTISTAGE OPTIMISTIC TD METHODS TOGETHER WITH
EXPECTIMAX SEARCH IN THE 4×6 -TUPLE AND THE 8×6 -TUPLE NETWORKS

Network	Search	Average Score	8192 [%]	16384 [%]	32768 [%]
2-stage 4×6 -tuple ($V_{\text{init}}=0$)	1-ply	291 428 \pm 5229	93.30 \pm 1.21%	68.58 \pm 1.05%	7.72 \pm 1.80%
	3-ply	463 784 \pm 12 569	99.72 \pm 0.15%	96.16 \pm 0.36%	46.97 \pm 1.28%
	5-ply	478 831 \pm 13 174	100.00 \pm 0.00%	98.40 \pm 1.10%	53.80 \pm 3.29%
2-stage 4×6 -tuple ($V_{\text{init}}=320k$)	1-ply	290 549 \pm 4105	93.03 \pm 0.62%	67.88 \pm 1.36%	8.42 \pm 0.44%
	3-ply	474 167 \pm 7087	99.70 \pm 0.15%	95.41 \pm 0.98%	46.02 \pm 2.53%
	5-ply	485 588 \pm 30 802	100.00 \pm 0.00%	97.00 \pm 4.00%	53.00 \pm 9.70%
2-stage 8×6 -tuple ($V_{\text{init}}=0$)	1-ply	386 078 \pm 17 104	97.37 \pm 0.23%	85.25 \pm 1.23%	26.67 \pm 1.64%
	3-ply	506 486 \pm 50 356	99.66 \pm 0.11%	96.89 \pm 0.32%	55.50 \pm 3.07%
	5-ply	516 761 \pm 67 952	100.00 \pm 0.00%	97.40 \pm 3.35%	56.80 \pm 8.53%
2-stage 8×6 -tuple ($V_{\text{init}}=320k$)	1-ply	404 288 \pm 2583	97.25 \pm 0.11%	85.37 \pm 0.39%	30.17 \pm 0.48%
	3-ply	538 582 \pm 5692	99.70 \pm 0.12%	96.83 \pm 0.32%	57.58 \pm 0.81%
	5-ply	581 896 \pm 14 765	99.80 \pm 0.89%	98.60 \pm 2.68%	66.40 \pm 6.42%
3-stage 8×6 -tuple ($V_{\text{init}}=320k$)	1-ply	412 492 \pm 3666	97.23 \pm 0.13%	85.34 \pm 0.35%	33.55 \pm 0.91%
	3-ply	545 231 \pm 6479	99.66 \pm 0.12%	96.90 \pm 0.44%	60.48 \pm 0.19%
	5-ply	588 426 \pm 30 723	99.80 \pm 0.89%	98.20 \pm 2.19%	72.80 \pm 8.17%

The 4×6 -tuple and the 8×6 -tuple were trained with OTC and OTD+TC, respectively.




Fig. 10. States before and after tile-downgrading. The original state (a) is translated to its downgraded state (b), where the highlighted tiles indicate that the values have been halved by tile-downgrading.

using their appropriate optimistic methods and hyperparameters, V_{init} and P_{TC} , as mentioned in previous experiment.

Table VIII gives the performance of the 2-stage 4×6 -tuple and 8×6 -tuple network with 1-ply, 3-ply, and 5-ply expectimax search. For comparison, the nonoptimistic 2-stage networks are also provided as baselines. In the second stage, average scores were improved by around $10\% \pm 5\%$. For further improvement, we attempted to add the third stage to the 2-stage 8×6 -tuple network by starting with 16384-tile + 8192-tile. However, such a 3-stage 8×6 -tuple network only improved the average scores by less than 2%, as given in Table VIII. Our conjecture is that the hyperparameters V_{init} and P_{TC} for the 8×6 -tuple, tuned for the first stage, did not fit the third stage.

3) Tile-Downgrading: Tile-downgrading expectimax search is a technique that searches states by translating them into downgraded states. Let us illustrate it by an example as in Fig. 10, where (a) is an original state and (b) is its downgraded state. In the original state (a), 256-tile is the largest missing tile not exceeding the largest tile, i.e., 32768-tile. This technique is to halve all tile values larger than 256. Thus, the downgraded state is derived as (b). Then, we estimate the value of the downgraded state, instead of the original state. Since states with large tiles have relatively less chance to be trained, their state estimations are relatively less accurate. For the downgraded states, their tile values are relatively smaller, implying higher chances to be trained and higher accuracy.

The tile-downgrading can only be applied to the root state of a search tree when the following two conditions hold in the root: there exists a missing tile like 256-tile as above; and there exists 32768-tiles or larger tiles, which is so chosen since it performed best in most of our experiments. Let a root state s satisfy the two conditions. The tile-downgrading expectimax search is to translate s into a downgraded state \tilde{s} and then use the expectimax search to choose the best action \tilde{a} to play. Note that no more tile-downgrading is required inside the search. Since this technique preserves the puzzle structure while halving some tile values, the action \tilde{a} of \tilde{s} can be directly applied to the original state s .

The tile-downgrading expectimax search was evaluated on the 2-stage 8×6 -tuple network trained with $V_{\text{init}} = 320\,000$. From Table IX, the tile-downgrading method improved the average score to 625 377, and the 32768-tile reaching rate to 72% with 6-ply search. Interestingly, this technique also significantly improved the chance of creating 65536-tiles to roughly 0.02% with 3-ply search, which is the highest 65536-tile reaching rate. Note that 65536-tiles were not reached for 5-ply and 6-ply searches since insufficient episodes (100 only) were evaluated for 5-ply and 6-ply.

D. Comparison to SOTA

In order to compare with the previous SOTA [5], [39], we used the same network, namely the 5×6 -tuple network (see Fig. 2) with 2-stage OTC with downgrading, and additionally we used a larger network, the 8×6 -tuple network (see Fig. 3) with 2-stage OTD+TC with downgrading. We tested 1-ply, 3-ply, and 5-ply search for these networks, and the results of the best n -tuple network of each method are given in Table IX. From Table IX, both our methods outperformed the 16-stage 5×6 -tuple network with TC in [5] in terms of fixed-depth search.

Furthermore, for the 2-stage 8×6 -tuple network, we tested an additional 6-ply search with 100 episodes. The average search speed was about 2.5 moves/s by using a single thread of an Intel Core i9-7960X. The 2-stage OTD+TC 8×6 -tuple network with 6-ply tile-downgrading expectimax search achieved an average score of 625 377, and a rate of 72% to reach 32768-tiles, which is also superior to those with the 1000-ms search in [5]. Table X gives the improvements of each technique from the nonoptimistic method (V0) to the best design (V5) in the upper six lines, and ablation studies (X1 to X4) in the lower four lines. To our knowledge, the results are SOTA in terms of average score and the reaching rates of large tiles. Although we used 3 more 6-tuples, we only required 2-stage when compared to their work with 16-stage. Since each extra stage requires one more set of tuples, we use far less memory in total, namely 16 (2×8) 6-tuples, versus 80 (16×5) 6-tuples in [5].

However, compared with nonoptimistic methods, our OI methods require more computing resources to explore large state spaces. The previous SOTA result [5] was trained with 4×10^{10} actions, which corresponds to approximately 40 million episodes or less. In contrast, our SOTA result was trained with 200 million episodes. Nevertheless, such an amount of training is still affordable. It took us 66 h to complete when using a C++ implementation [40] on an Intel E5-2698 v4 processor.

TABLE IX
COMPARISON BETWEEN SOTA OPTIMISTIC METHODS AND THE PREVIOUS SOTA METHOD

Authors	Methods	Weights	Search	Average Score	8192 [%]	16 384 [%]	32 768 [%]	# Games
Jaśkowski [5]	MS-TC 16-stage 5×6-tuple network	1342.2M	1-ply	324 710 ± 11 043	90%	68%	19%	1000
			3-ply	511 759 ± 12 021	99%	92%	50%	1000
			5-ply	545 833 ± 21 500	100%	97%	54%	300
			1000ms	609 104 ± 38 433	98%	97%	70%	100
This work	MS-OTC 2-stage 5×6-tuple network with tile-downgrading	167.8M	1-ply	331 073 ± 6699	94.22 ± 0.12%	74.41 ± 0.44%	15.38 ± 0.62%	1 000 000
			3-ply	526 178 ± 21 752	99.69 ± 0.18%	95.89 ± 0.77%	51.07 ± 3.87%	10 000
			5-ply	574 245 ± 25 727	99.80 ± 0.89%	97.00 ± 2.00%	60.60 ± 5.02%	100
	MS-OTD+TC 2-stage 8×6-tuple network with tile-downgrading	268.4M	1-ply	412 785 ± 2208	97.24 ± 0.12%	85.39 ± 0.35%	30.16 ± 0.38%	1 000 000
			3-ply	563 316 ± 5650	99.63 ± 0.19%	96.88 ± 0.12%	57.90 ± 1.63%	10 000
			5-ply	608 679 ± 42 177	99.80 ± 0.89%	97.80 ± 1.67%	67.40 ± 12.21%	100
			6-ply	625 377 ± 40 936	99.80 ± 0.89%	98.80 ± 0.89%	72.00 ± 12.00%	100

TABLE X
IMPROVEMENT OF INDIVIDUAL TECHNIQUES TO THE BEST RESULT

Method	Average Score	32 768 [%]
V0 = 1-stage 8×6-tuple TD ($V_{\text{init}}=0$)	309 208 ± 5788	0.00 ± 0.00%
V1 = V0 + OI ($V_{\text{init}}=320k$)	361 471 ± 5473	21.75 ± 0.74%
V2 = V1 + TC fine-tuning ($P_{\text{TC}}=10\%$)	370 194 ± 4366	22.18 ± 0.92%
V3 = V2 + multistage (2-stage)	404 288 ± 2583	30.17 ± 0.48%
V4 = V3 + expectimax search (6-ply)	586 583 ± 17 043	65.40 ± 2.28%
V5 = V4 + tile-downgrading	625 377 ± 40 936	72.00 ± 12.00%
X1 = V5 - OI	592 390 ± 39 073	63.60 ± 10.35%
X2 = V5 - TC fine-tuning	574 779 ± 22 280	59.80 ± 7.40%
X3 = V5 - multistage	574 150 ± 29 315	59.80 ± 10.14%
X4 = V5 - expectimax search	412 785 ± 2208	30.16 ± 0.38%

V0–V5 present the incremental improvements to the best design (V5).

X1–X4 present the ablation studies of the best design.

V. DISCUSSION

In this section, we discuss some other techniques related to 2048, including weight promotion [5], [26], redundant encoding [5], and carousel shaping [5].

Weight promotion can be used to speed up the multistage training by allowing the weights of n -tuple networks to be copied to the corresponding weights in the next stage upon its first access. This method is similar to OI in the sense that many weights are initialized higher than zeros. However, when compared to OI, their method only copies weights with small tiles, and those weights with large tiles, say 16384-tile in the next stage, are still initialized to zeros. In OI, these weights are large numbers leading to more exploration.

Redundant encoding improves the n -tuple network by adding some sub-tuples. OI may be applicable to this technique with an initialization scheme different from that in (12). Carousel shaping ensures each stage receives the same training amount. Since this article focuses on the impact of OI, further research on incorporating these methods into OI is left open.

VI. CONCLUSION

For the issue of exploration on 2048, we propose optimistic TD learning to improve the performance. Our approach significantly improves the learning quality of n -tuple networks. The significance of this article is summarized as follows.

First, we improve the common TD methods with OI by using a hyperparameter V_{init} to initialize the network weights, and demonstrate that OTD and OTC learning with $V_{\text{init}} = 320\,000$ significantly improve the performance, especially the chance of obtaining 32768-tiles, as shown in Section IV-A. Second, we

propose OTD+TC learning, a hybrid learning paradigm that combines the advantages of both TD and TC for encouraging exploration and exploitation, respectively. A hyperparameter P_{TC} is used to control the proportion of TC fine-tuning phase. We observe that OTD+TC with $V_{\text{init}} = 320\,000$ and $P_{\text{TC}} = 10\%$ outperforms both OTD and OTC learning for larger networks, as described in Section IV-B. Third, we show that we do not need as many stages as in [5], thereby significantly reducing the required network weights, since OI effectively improves the learning quality as given in Section IV-C2.

Furthermore, we use tile-downgrading to improve the search quality as shown in Section IV-C3, and the results in Section IV-D shown to outperform those in the previous SOTA [5] in terms of average score and 32768-tile reaching rate. Namely, the design with a 2-stage 8×6 -tuple network trained by OTD+TC with $V_{\text{init}} = 320\,000$ and $P_{\text{TC}} = 10\%$ achieved an average score of 625 377 and a rate of 72% reaching 32768-tiles through 6-ply tile-downgrading expectimax search. In addition, for sufficiently large tests, 65536-tiles are reached at a rate of 0.02%.

We conclude that for 2048 and similar games, learning with explicit exploration is useful even if the environment seems to be stochastic enough. This article demonstrates that the optimistic method is promising for stochastic games.

REFERENCES

- [1] G. Cirulli, “2048, success and me,” May–Sep., 2014. Accessed: Jun. 16, 2020. [Online]. Available: <https://web.archive.org/web/20140517/http://gabrielecirulli.com/articles/2048-success-and-me>
- [2] M. Szubert and W. Jaśkowski, “Temporal difference learning of ntuple networks for the game 2048,” in *Proc. IEEE Conf. Comput. Intell. Games*, 2014, pp. 1–8, doi: [10.1109/CIG.2014.693290](https://doi.org/10.1109/CIG.2014.693290). [Online]. Available: <http://www.cs.put.poznan.pl/mszubert/pub/szubert2014cig.pdf>
- [3] K.-H. Yeh, I.-C. Wu, C.-H. Hsueh, C.-C. Chang, C.-C. Liang, and H. Chiang, “Multistage temporal difference learning for 2048-like games,” *IEEE Trans. Comput. Intell. AI Games*, vol. 9, no. 4, pp. 369–380, Dec. 2017.
- [4] R. Mehta, “2048 is (PSPACE) hard, but sometimes easy,” Electron. Colloq. Comput. Complex., Tech. Rep. TR14-116, 2014. [Online]. Available: <https://eccc.weizmann.ac.il/report/2014/116/download>
- [5] W. Jaśkowski, “Mastering 2048 with delayed temporal coherence learning, multistage weight promotion, redundant encoding and carousel shaping,” *IEEE Trans. Games*, vol. 10, no. 1, pp. 3–14, Mar. 2018.
- [6] H. Guei, T.-H. Wei, and I.-C. Wu, “2048-like games for teaching reinforcement learning,” *ICGA J.*, vol. 42, no. 1, pp. 14–37, May 28, 2020.
- [7] W. Konen, “General board game playing for education and research in generic AI game learning,” in *Proc. IEEE Conf. Games*, 2019, pp. 1–8, doi: [10.1109/CIG.2019.8848070](https://doi.org/10.1109/CIG.2019.8848070).
- [8] K. Matsuzaki, “Systematic selection of ntuple networks with consideration of interinfluence for game 2048,” in *Proc. 21st Int. Conf. Technol. Appl. Artif. Intell.*, 2016, pp. 186–193.
- [9] R. S. Sutton and A. G. Barto, *Reinforcement Learning: An Introduction*, 1st ed. Cambridge, MA, USA: MIT Press, 1998. [Online]. Available: <http://incompleteideas.net/book/first/ebook/the-book.html>

- [10] I.-C. Wu, K.-H. Yeh, C.-C. Liang, C.-C. Chang, and H. Chiang, "Multistage temporal difference learning for 2048," in *Proc. 19th Int. Conf. Technol. Appl. Artif. Intell.*, 2014, pp. 366–378.
- [11] K. Oka and K. Matsuzaki, "Systematic selection of ntuple networks for 2048," in *Proc. 9th Int. Conf. Comput. Games*, 2016, pp. 81–92.
- [12] D. F. Beal and M. C. Smith, "Temporal coherence and prediction decay in TD learning," in *Proc. 16th Int. Joint Conf. Artif. Intell.*, 1999, vol. 1, pp. 564–569. [Online]. Available: <http://ijcai.org/Proceedings/99-1/Papers/081.pdf>
- [13] M. C. Machado, S. Srinivasan, and M. Bowling, "Domain-independent optimistic initialization for reinforcement learning," in *Proc. 29th AAAI Conf. Artif. Intell. Workshop Learn. Gen. Competition Video Games*, 2015, pp. 29–31. [Online]. Available: <https://aaai.org/ocs/index.php/WS/AAAIW15/paper/view/10103/10218>
- [14] G. Tesauro, "TD-Gammon, a self-teaching backgammon program, achieves master-level play," *Neural Comput.*, vol. 6, no. 2, pp. 215–219, Mar. 1994.
- [15] J. Schaeffer, M. Hlynka, and V. Jussila, "Temporal difference learning applied to a high-performance game-playing program," in *Proc. 17th Int. Joint Conf. Artif. Intell.*, 2001, pp. 529–534.
- [16] J. Baxter, A. Tridgell, and L. Weaver, "Learning to play chess using temporal differences," *Mach. Learn.*, vol. 40, no. 3, pp. 243–263, Sep. 2000.
- [17] D. F. Beal and M. C. Smith, "First results from using temporal difference learning in Shogi," in *Proc. 1st Int. Conf. Comput. Games*, 1998, pp. 113–125.
- [18] D. Silver, "Reinforcement learning and simulation-based search in computer go," Ph.D. dissertation, Dept. Comput. Sci., Univ. Alberta, Edmonton, AB, Canada, 2009.
- [19] I.-C. Wu, H.-T. Tsai, H.-H. Lin, Y.-S. Lin, C.-M. Chang, and P.-H. Lin, "Temporal difference learning for connect6," in *Proc. 13th Int. Conf. Adv. Comput. Games*, 2011, pp. 121–133.
- [20] M. Thill, P. Koch, and W. Konen, "Reinforcement learning with N-tuples on the game connect-4," in *Proc. 12th Int. Conf. Parallel Problem Solving Nature*, 2012, pp. 184–194.
- [21] S. M. Lucas, "Learning to play othello with N-tuple systems," *Australian J. Intell. Inf. Process.*, vol. 9, no. 4, pp. 1–20, Feb. 2008. [Online]. Available: <http://ajipps.com.au/papers/V9.4/V9N4.1%20-%20Learning%20to%20Play%20Othello%20with%20Ntuple%20Systems.pdf>
- [22] R. S. Sutton, "Learning to predict by the methods of temporal differences," *Mach. Learn.*, vol. 3, no. 1, pp. 9–44, Aug. 1988.
- [23] K. Matsuzaki, "Developing a 2048 player with backward temporal coherence learning and restart," in *Proc. 15th Int. Conf. Adv. Comput. Games*, 2017, pp. 176–187.
- [24] H. Guei, T.-H. Wei, and I.-C. Wu, "Using 2048-like games as a pedagogical tool for reinforcement learning," in *Proc. 10th Int. Conf. Comput. Games*, Mar. 2019, pp. 281–293.
- [25] A. G. Barto and S. Mahadevan, "Recent advances in hierarchical reinforcement learning," *Discrete Event Dyn. Syst.*, vol. 13, no. 1/2, pp. 41–77, Oct. 2003.
- [26] K. Matsuzaki, "Evaluation of multi-staging and weight promotion for game 2048," Kochi Univ. Technol., Kami, Japan, Tech. Rep. KUTBTR2017, Oct. 27, 2017. [Online]. Available: <http://hdl.handle.net/10173/1564>
- [27] W. W. Bledsoe and I. Browning, "Pattern recognition and reading by machine," in *Proc. Eastern Joint Comput. Conf.*, 1959, pp. 225–232. [Online]. Available: <https://dl.acm.org/doi/pdf/10.1145/1460299.1460326>
- [28] B. W. Ballard, "The *-minimax search procedure for trees containing chance nodes," *Artif. Intell.*, vol. 21, no. 3, pp. 327–350, Sep. 1983.
- [29] E. Melkó and B. Nagy, "Optimal strategy in games with chance nodes," *Acta Cybern.*, vol. 18, no. 2, pp. 171–192, Jan. 2007.
- [30] N. Pezzotti, "An artificial intelligence for the 2048 game," Mar. 26, 2014. Accessed: Jun. 16, 2020. [Online]. Available: <https://diaryofatinker.blogspot.com/2014/03/an-artificial-intelligence-for-2048-game.html>
- [31] A. L. Zobrist, "A new hashing method with application for game playing," *ICGA J.*, vol. 13, no. 2, pp. 69–73, Jun. 1990.
- [32] K.-H. Yeh, "2048 AI." Accessed: Jun. 16, 2020. [Online]. Available: <https://github.com/tnmichael309/2048AI>
- [33] A. Appleby, "MurmurHash3 64-bit finalizer," Aug. 1, 2016. Accessed: Jun. 16, 2020. [Online]. Available: <https://github.com/kcwu/2048-c>
- [34] K.-C. Wu, "2048-c." Accessed: Jun. 16, 2020. [Online]. Available: <https://github.com/kcwu/2048-c>
- [35] I. Szita and A. Lorincz, "The many faces of optimism: A unifying approach," in *Proc. 25th Int. Conf. Mach. Learn.*, 2008, pp. 1048–1055. [Online]. Available: <https://dl.acm.org/doi/pdf/10.1145/1390156.1390288>
- [36] V. Kuleshov and D. Precup, "Algorithms for multi-armed bandit problems," *J. Mach. Learn. Res.*, vol. 1, pp. 1–48, 2000. [Online]. Available: <https://arxiv.org/pdf/1402.6028.pdf>
- [37] E. Even-Dar and Y. Mansour, "Convergence of optimistic and incremental Q-learning," in *Proc. 14th Int. Conf. Neural Inf. Process. Syst.*, 2001, pp. 1499–1506. [Online]. Available: <http://papers.nips.cc/paper/1944-convergence-of-optimistic-and-incremental-q-learning.pdf>
- [38] N. K. Jong, T. Hester, and P. Stone, "The utility of temporal abstraction in reinforcement learning," in *Proc. 7th Int. Joint Conf. Auton. Agents Multiagent Syst.*, 2008, pp. 299–306. [Online]. Available: <https://dl.acm.org/doi/pdf/10.5555/1402383.1402429>
- [39] W. Jaśkowski and A. Szczepański, "2048 AI." Accessed: Jun. 16, 2020. [Online]. Available: <https://github.com/aszczepanski/2048>
- [40] H. Guei, "moporgic/TDL2048+." Accessed: Apr. 8, 2021. [Online]. Available: <https://github.com/moporgic/TDL2048>



Hung Guei is currently working toward the Ph.D. degree in computer science with the Department of Computer Science, National Yang Ming Chiao Tung University, Hsinchu, Taiwan.

His research interests include artificial intelligence, machine learning, computer games, and grid computing.



Lung-Pin Chen (Member, IEEE) received the M.S. degree from the National Chung-Cheng University, Chiayi, Taiwan, in 1993, and the Ph.D. degree from the National Yang Ming Chiao Tung University, Hsinchu, Taiwan, in 1999, both in computer science.

He is currently an Associate Professor with the Department of Computer Science and Information Engineering, Tunghai University, Taichung, Taiwan. His research interests include distributed algorithm, cloud computing, computational intelligence, and machine learning.



I-Chen Wu (Senior Member, IEEE) received the B.S. degree in electronic engineering and the M.S. degree in computer science from National Taiwan University, Taipei, Taiwan, in 1982 and 1984, respectively, and the Ph.D. degree in computer science from Carnegie-Mellon University, Pittsburgh, PA, USA, in 1993.

He is currently the Executive Officer of Artificial Intelligence Computing Center at Academia Sinica, a Research Fellow of Research Center for IT Innovation at Academia Sinica, and also a Professor with the

Department of Computer Science, National Yang Ming Chiao Tung University, Hsinchu, Taiwan. His research achievements include several state-of-the-art game playing programs, such as CGI for Go and Chimo for Chinese Chess, winning more than 30 gold medals in international tournaments, like Computer Olympiad. He wrote more than 150 technical papers, and served as chairs and member of committees in more than 30 academic conferences and organizations, including the conference chair of the 2015 IEEE CIG Conference. His research interests include computer games and deep reinforcement learning.

Dr. Wu is the Editor-in-Chief of *ICGA Journal* and an Associate Editor of the *IEEE TRANSACTIONS ON GAMES*. He is currently the Vice President of the International Computer Games Association and the President of the Taiwanese Computer Games Association; and was the President of the Taiwanese Association for Artificial Intelligence in 2015–2017.